

Use of AI for Log Analysis in CI/CD Pipelines

Bachelor Thesis - Defence

Maid Ališić

July 23, 2025

FH Oberösterreich · Campus Hagenberg

Supervisor: FH-Prof. Dipl.-Ing. Dr. Stefan Wagner

Road map

Research questions

Problem context

Method

Architecture

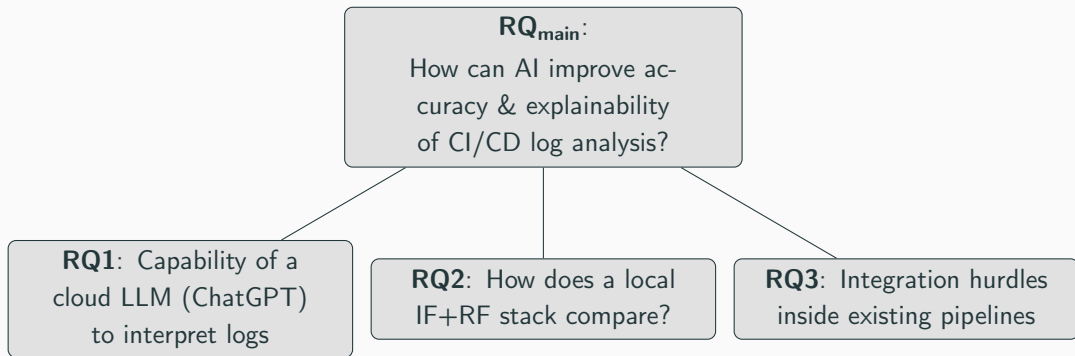
Data & evaluation

Results

Impact

Research questions

Research questions



Problem context

Why do we care?

- CI/CD emits \approx 10-20 GB of build, test & deploy logs *per day*.

Why do we care?

- CI/CD emits \approx 10-20 GB of build, test & deploy logs *per day*.
- Manual grep slows the merge queue; critical faults slip through.

Why do we care?

- CI/CD emits \approx 10-20 GB of build, test & deploy logs *per day*.
- Manual grep slows the merge queue; critical faults slip through.
- Business **Service-Level Objective**: feedback within \leq 200 ms per pipeline.

Why do we care?

- CI/CD emits \approx 10-20 GB of build, test & deploy logs *per day*.
- Manual grep slows the merge queue; critical faults slip through.
- Business **Service-Level Objective**: feedback within \leq 200 ms per pipeline.
- Logs may expose customer IDs, therefore they must remain on-premises (no cloud export).

1. **Context-sensitivity** - identical tokens can be harmless or fatal.

1. **Context-sensitivity** - identical tokens can be harmless or fatal.
2. **Concept drift** - each merge may rename tests or switches.

1. **Context-sensitivity** - identical tokens can be harmless or fatal.
2. **Concept drift** - each merge may rename tests or switches.
3. **Latency pressure** - analysis must finish before runner teardown.

Operational pain points

1. **Context-sensitivity** - identical tokens can be harmless or fatal.
2. **Concept drift** - each merge may rename tests or switches.
3. **Latency pressure** - analysis must finish before runner teardown.
4. **Alert fatigue** - regex rule sets grow without bound.

Method

1. **Normalise** - strip timestamps, colours, IDs.



Vectorisation pipeline ①

1. **Normalise** - strip timestamps, colours, IDs.
2. **Tokenise** into uni- and bi-grams.



Vectorisation pipeline ①

1. **Normalise** - strip timestamps, colours, IDs.
2. **Tokenise** into uni- and bi-grams.
3. Weight with TF-IDF.




Vectorisation pipeline ①

1. **Normalise** - strip timestamps, colours, IDs.
2. **Tokenise** into uni- and bi-grams.
3. Weight with TF-IDF.
4. Produce sparse vector.




Isolation Forest ② - intuition

- Random binary partitioning isolates unusual lines in fewer splits.




Isolation Forest ② - intuition

- Random binary partitioning isolates unusual lines in fewer splits.
- Score $s(x) = 2^{-h(x)/c(n)} \in [0,1]$
if high \rightarrow outlier.




Isolation Forest ② - intuition

- Random binary partitioning isolates unusual lines in fewer splits.
- Score $s(x) = 2^{-h(x)/c(n)} \in [0,1]$ if high \rightarrow outlier.
- CPU-only: $\approx 30 \mu s$ per line.




Random Forest ③ - error labelling

- Maps each flagged line to a domain-specific error category.




Random Forest ③ - error labelling

- Maps each flagged line to a domain-specific error category.
- Majority vote = deterministic, auditable output.



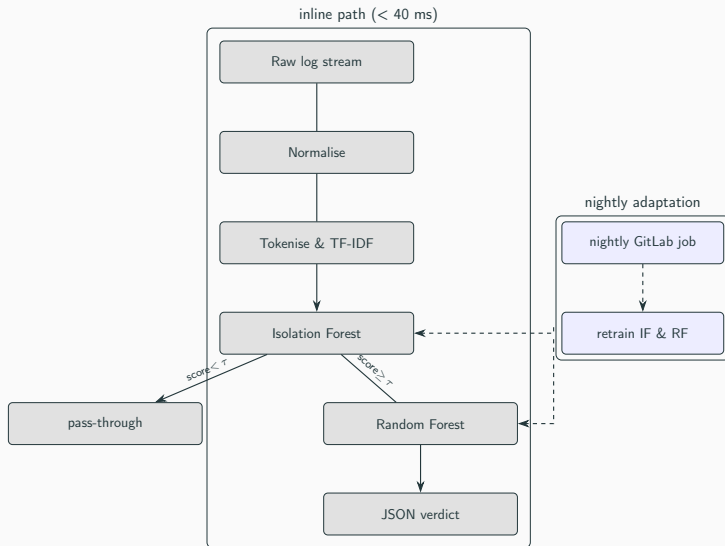
Random Forest ③ - error labelling

- Maps each flagged line to a domain-specific error category.
- Majority vote = deterministic, auditable output.
- Nightly retrain < 90 s; warm-start handles drift.




Architecture

End-to-end pipeline (< 40 ms inline)




Data & evaluation

- 117 k *macOS* logs + 655 k *OpenSSH* logs




Datasets & metrics

- 117 k *macOS* logs + 655 k *OpenSSH* logs
- 504 labelled anomalies (class imbalance \approx 1 : 200)




- 117 k *macOS* logs + 655 k *OpenSSH* logs
- 504 labelled anomalies (class imbalance \approx 1 : 200)
- Split 70 / 15 / 15 % (training / validation / testing)



Datasets & metrics

- 117 k *macOS* logs + 655 k *OpenSSH* logs
- 504 labelled anomalies (class imbalance $\approx 1 : 200$)
- Split 70 / 15 / 15 % (training / validation / testing)
- Metrics: Precision, Recall and F_1



Results

Headline numbers

	Precision	Recall	F_1
Detection (Isolation Forest)	0.91	0.88	0.89
Classification (Random Forest)	0.99	0.99	0.99
Regex Baseline	0.286	0.286	0.286

$$F_1 = 2 \cdot \frac{P \cdot R}{P + R}$$

Impact

- **Latency:** minutes → **milliseconds** (inline verdict).

- **Latency:** minutes → **milliseconds** (inline verdict).
- **Cost-free:** 2.3 k lines of code, CPU-only, no token fees.

- **Latency:** minutes → **milliseconds** (inline verdict).
- **Cost-free:** 2.3 k lines of code, CPU-only, no token fees.
- **GDPR compliant:** logs never leave the VPN.

Wrap-up

Light-weight on-prem ML matches AIOps SaaS
without latency, cost or privacy pain.

Questions welcome - thank you!